

# Data Visualization in R with ggplot2

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Acknowledgments</b>	<b>2</b>
<b>3</b>	<b>Resources</b>	<b>2</b>
<b>4</b>	<b>Getting Started</b>	<b>2</b>
4.1	Prerequisites . . . . .	2
4.2	Software Requirements . . . . .	3
4.3	Installing ggplot2 . . . . .	3
4.4	Installing ggplot2 Extentions . . . . .	3
<b>5</b>	<b>Grammer of Graphics</b>	<b>4</b>
5.1	Building Blocks of a Graph . . . . .	5
<b>6</b>	<b>Basic Plots</b>	<b>5</b>
6.1	Loading ggplot2 . . . . .	5
6.2	Dataset . . . . .	6
6.3	Scatter Plot . . . . .	6
6.4	Exercise . . . . .	8
<b>7</b>	<b>Geoms and Statistics</b>	<b>10</b>
7.1	Exercise . . . . .	16
<b>8</b>	<b>Scales</b>	<b>19</b>
<b>9</b>	<b>Text and Labels</b>	<b>22</b>
<b>10</b>	<b>Smoother</b>	<b>26</b>
<b>11</b>	<b>Theme and Title</b>	<b>31</b>
<b>12</b>	<b>Facets</b>	<b>41</b>
<b>13</b>	<b>Challenge</b>	<b>44</b>
13.1	Recreating the Economist Graph . . . . .	44

# 1 Introduction

This course introduces you to data visualization in R using the ggplot2 package. The ggplot2 package implements the grammar of graphics concepts for creating visually appealing and professional looking graphics in R. Basic knowledge of working with datasets in R is essential but experience with plotting functions is not required.

By the end of the course you will be able to:

- Create scatterplots, histograms, line graphs, and boxplots
- Add chart labels, axis labels and legends to the plots
- Apply statistical transformations to the plots
- Change various attributes of plot layers including color, shape, size and scale
- Create mutiple plots in a single figure using facets
- Apply themes to change the appearance of the plots

---

*Last Updated: Oct 19, 2017 12:41 AM*

## 2 Acknowledgments

Content of this workshop is based on the following:

- ggplot2 tutorial from Harvard University
- ggplot2 Workshop (Vanderbilt, 2007)

This work is licensed under a Creative Commons Attribution-ShareAlike 4.0 International License.

## 3 Resources

- Google
- Data Visualization with ggplot2 Cheat Sheet
- ggplot2 Documentation
- R Graphics Cookbook
- ggplot2 tutorial from Harvard University
- ggplot2 Workshop (Vanderbilt, 2007)

## 4 Getting Started

### 4.1 Prerequisites

Basic knowledge of working with datasets in R is essential. This course assumes that you're comfortable with reading and manipulating datasets, working with script files, and navigating in RStudio. Experience with plotting functions in R is helpful but not required.

## 4.2 Software Requirements

### 4.2.1 R and RStudio

Recent versions of R (version 3.2 or newer) and RStudio (version 0.99 or above) are required.

You can download the latest versions from the links below:

- [Download R](#)
- [Download RStudio](#)

You can find out the version of R installed by typing `version` at the console:

```
version

##
## platform      x86_64-pc-linux-gnu
## arch          x86_64
## os            linux-gnu
## system        x86_64, linux-gnu
## status
## major         3
## minor         4.2
## year          2017
## month         01
## day           27
## svn rev       73369
## language      R
## version.string R version 3.4.2 (2017-01-27)
## nickname      Short Summer
```

## 4.3 Installing ggplot2

If you don't have ggplot2 installed, you can install it using the `install.packages()` function:

```
install.packages("ggplot2")
```

You can find out the version of ggplot installed using the `packageVersion()` function:

```
packageVersion("ggplot2")
```


```
## [1] '2.2.1'
```

## 4.4 Installing ggplot2 Extensions

We need the following ggplot extensions for this tutorial:

```
install.packages("scales")
install.packages("ggrepel")
install.packages("ggthemes")
```

## 5 Grammer of Graphics



Wilkinson, L. (2006). *The grammar of graphics*. Springer Science & Business Media.

---

## 5.1 Building Blocks of a Graph

- Data
- Aesthetic mapping
- Geometric object
- Statistical transformations
- Scales
- Coordinate system
- Position adjustments
- Faceting




Figure 1:

## 6 Basic Plots

### 6.1 Loading ggplot2

Like any other R package, you must load **ggplot2** using the **library** function before you can use any of the functionality that it offers. We will also load the extensions that provide additional features:

```
library(ggplot2)
library(ggrepel)
library(ggthemes)
library(scales)
```

## 6.2 Dataset

Let's start by loading the housing dataset:

```
housing <- read.csv("https://raw.githubusercontent.com/altaf-ali/ggplot_tutorial/master/data/housing.csv")
```

Now, let's see what the dataset looks like:

```
head(housing)
```

```
##   State Region      Date Home.Value Structure.Cost Land.Value
## 1    AK  West 2010-03-01   224952         160599    64352
## 2    AK  West 2010-06-01   225511         160252    65259
## 3    AK  West 2009-09-01   225820         163791    62029
## 4    AK  West 2009-12-01   224994         161787    63207
## 5    AK  West 2007-12-01   234590         155400    79190
## 6    AK  West 2008-03-01   233714         157458    76256
##   Land.Share..Pct. Home.Price.Index Land.Price.Index Year Quarter
## 1                28.6             1.481             1.552 2010      1
## 2                28.9             1.484             1.576 2010      2
## 3                27.5             1.486             1.494 2009      3
## 4                28.1             1.481             1.524 2009      4
## 5                33.8             1.544             1.885 2007      4
## 6                32.6             1.538             1.817 2008      1
```

When dealing with date and time values, it's generally a good idea to convert them to the appropriate data type.

```
housing$Date <- as.Date(housing$Date)
```

Next, we create two subsets of the data, one with housing prices only from New York, and another one with housing prices from 9 states in the North East.

```
newyork <- subset(housing, State == "NY")
northeast <- subset(housing, Region == "N. East")
```


## 6.3 Scatter Plot

Now we're ready to plot. Everything starts with the `ggplot()` function which creates a plot object. The two arguments passed to `ggplot()` are:

Argument	Description
<code>data</code>	Dataset for the plot. It should be a data.frame or something that can be converted to data.frame
<code>mapping</code>	Aesthetic mappings for the plot

Using the `newyork` dataset, let's create a scatter plot with `Date` on the x-axis and `Home.Value` on the y-axis.

```
ggplot(newyork, aes(x = Date, y = Home.Value)) +  
  geom_point()
```



Now let's see which ggplot building blocks are active in the above example:

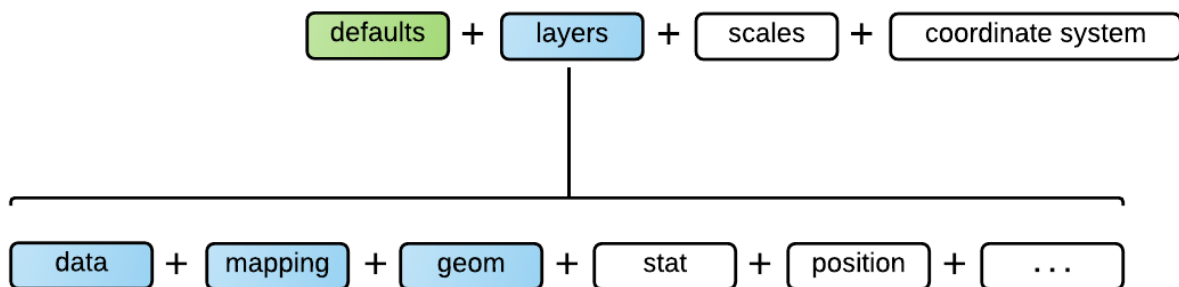



Figure 2:

Data	<code>newyork</code>
Mapping	<code>aes(x = Date, y = Home.Value)</code>
Geom	<code>geom_point()</code>

## 6.4 Exercise


Use the Data Visualization with ggplot2 Cheat Sheet or any other resource to find out how to complete the exercises.

1. Create a histogram of `Home.Value` using the `housing` data.




2. Create a box plot of `Home.Value` using `northeast` dataset with `State` on the x-axis






3. Create a line plot using `newyork` dataset with `Date` on the x-axis and `Home.Value` on the y-axis



4. Create a line plot using `northeast` dataset with `Date` on the x-axis and `Home.Value` on the y-axis and use a different color for each state



## 7 Geoms and Statistics

Geometric objects (geoms) define the basic shape of the elements on the plot.


- Every geom has a default statistic
- Every statistic has a default geom

You can get a list of all geoms using the online help in RStudio

```
help.search("geom_", package = "ggplot2")
```


Change the size of each bin:

```
ggplot(housing, aes(x = Home.Value)) +  
  geom_histogram(binwidth = 1000)
```



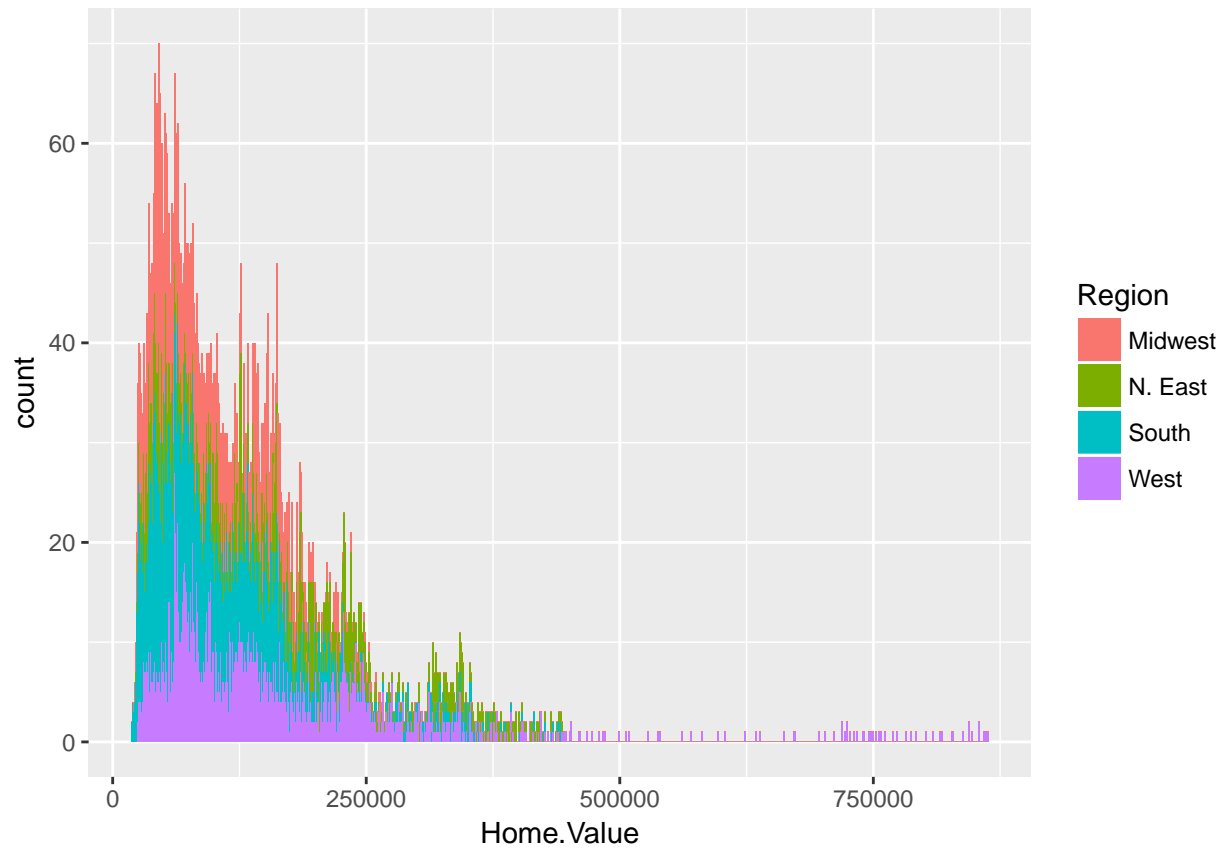
Add a mapping for the fill color:

```
ggplot(housing, aes(x = Home.Value, fill = Region)) +  
  geom_histogram(binwidth = 1000)
```




Mapping can also be specified in the geom:

```
ggplot(housing, aes(x = Home.Value)) +  
  geom_histogram(aes(fill = Region), binwidth = 1000)
```




Same plot can also be created using `stat_bin` transformation. The default geom for `stat_bin` is “area”

```
ggplot(housing, aes(x = Home.Value)) +  
  stat_bin(binwidth = 1000)
```




Change the default geom to "point"

```
ggplot(housing, aes(x = Home.Value)) +  
  stat_bin(geom = "point", binwidth = 1000)
```



Change the default geom to "line"

```
ggplot(housing, aes(x = Home.Value)) +  
  stat_bin(geom = "line", binwidth = 1000)
```




## 7.1 Exercise

Create a subset of housing data from New York since 2000

```
newyork2k <- subset(newyork, Year >= 2000)
```


1. Create a plot that includes multiple geometric objects, for example, lines and points.






2. Change the shape to be hollow diamond

HINT: Take a look at **Shape Scales** in the Data Visualization with ggplot2 Cheat Sheet



3. Change the size of the point based on **Land.Value** and color based on **Structure.Cost**




## 8 Scales

Let's create another subset that includes only the data from the first quarter of 2001.

```
housing2001q1 <- subset(housing, Year == 2001 & Quarter == 1)
```


And now we create a scatter plot with this dataset

```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +  
  geom_point()
```




Our dataset is skewed so in order to help with interpretation, let's change the x-axis to log scale

```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +  
  geom_point() +  
  scale_x_log10()
```




Now add a dollar sign in front of our axis labels

```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +  
  geom_point() +  
  scale_x_log10(labels = dollar) +  
  scale_y_continuous(labels = dollar)
```



Next we change the scale for the x-axis which is in a Date format and control the breaks for y-axis which is a continuous variable.

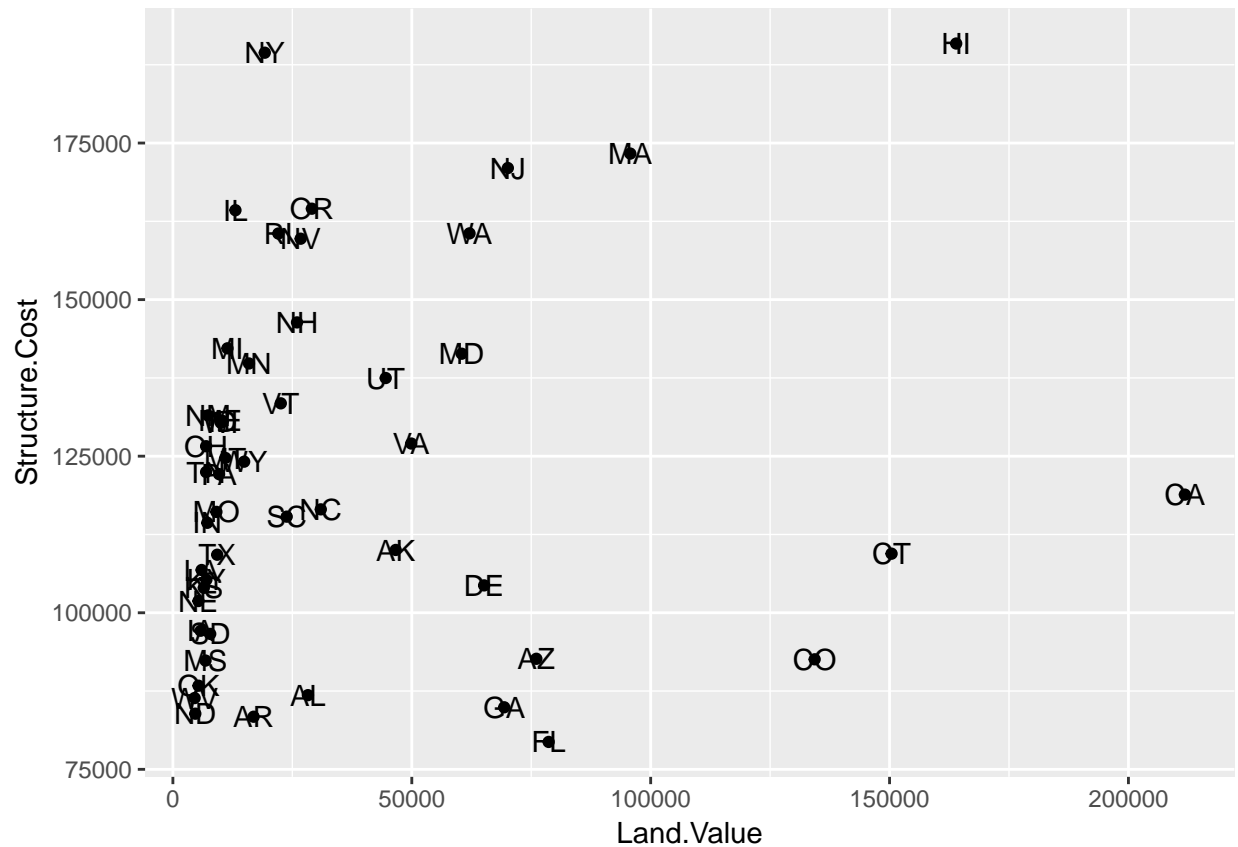
```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +
  geom_line() +
  scale_x_date(date_breaks = "3 year", date_minor_breaks = "1 year", date_labels = "%y") +
  scale_y_continuous(breaks = seq(0, 500000, 50000), labels = dollar)
```



## 9 Text and Labels


Let's continue with the subset of the data from the previous section and add text to the scatterplot.

```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +
  geom_point() +
  geom_text(aes(label = State))
```



The result isn't very nice as the labels overlap each other. Let's try the same with `geom_label()` instead which draws the text with a border around it.


```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +
  geom_point() +
  geom_label(aes(label = State))
```



The `ggrepel` extension we loaded earlier can also help fix this problem.


```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +  
  geom_point() +  
  geom_text_repel(aes(label = State))
```





And we can repel the labels with a border using `geom_label_repel()`.

```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +
  geom_point() +
  geom_label_repel(aes(label = State))
```




## 10 Smoother

Let's continue with the 2001 first quarter dataset and add a smoother.


```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +
  geom_point() +
  scale_x_log10() +
  stat_smooth()
```

```
## `geom_smooth()` using method = 'loess'
```




We can fit a linear model to our dataset

```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +  
  geom_point() +  
  scale_x_log10() +  
  stat_smooth(method = "lm")
```



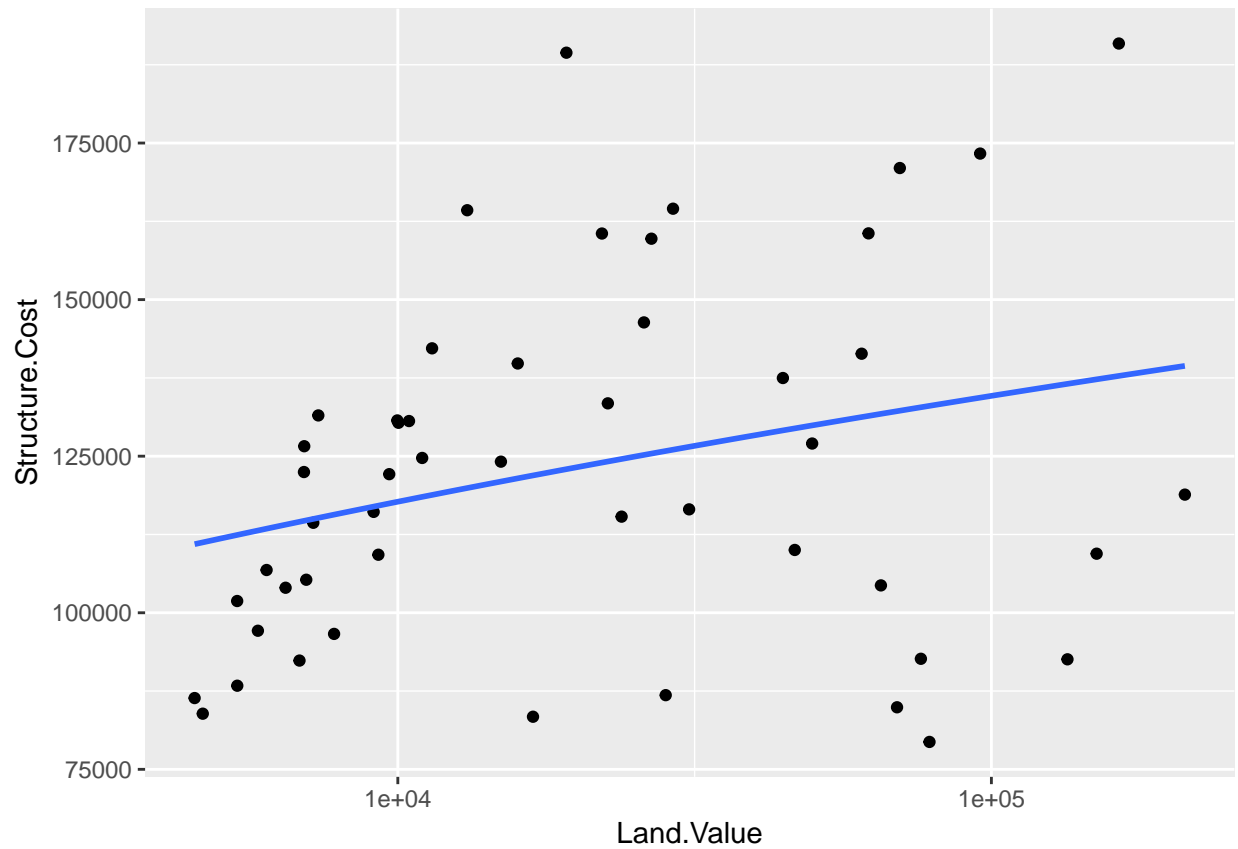
We can also specify the formula for the model

```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +  
  geom_point() +  
  scale_x_log10() +  
  stat_smooth(method = "lm", formula = y ~ log(x))
```




We can turn the turn off the confidence interval

```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +
  geom_point() +
  scale_x_log10() +
  stat_smooth(method = "lm", formula = y ~ log(x), se = FALSE)
```



Formula is specific to the type of model used. Here we're using a General Additive Model (GAM).


```
ggplot(housing2001q1, aes(x = Land.Value, y = Structure.Cost)) +  
  geom_point() +  
  scale_x_log10() +  
  stat_smooth(method = "gam", formula = y ~ s(x,k=10))
```



## 11 Theme and Title


First, let's try some of the themes from the `ggthemes` package

```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line() +  
  theme_stata()
```




```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line() +  
  theme_economist()
```







```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line() +  
  theme_wsj()
```



```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line() +  
  theme_solarized()
```




```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line() +  
  theme_fivethirtyeight()
```




We can also have complete control over the theme by customizing each element ourselves. Let's start with `theme_minimal()`

```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line() +  
  theme_minimal()
```




Now remove the minor grid lines

```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line() +  
  theme_minimal() +  
  theme(  
    panel.grid.minor = element_blank()  
  )
```




Next, we change the y-axis label

```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line() +  
  theme_minimal() +  
  theme(  
    panel.grid.minor = element_blank()  
  ) +  
  ylab("Home Value (US$)")
```



Then remove the x-axis title since the year is self explanatory


```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +
  geom_line() +
  theme_minimal() +
  theme(
    axis.title.x = element_blank(),
    panel.grid.minor = element_blank()
  ) +
  ylab("Home Value (US$)")
```



Finally, we can add a title to our plot

```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +
  geom_line() +
  theme_minimal() +
  theme(
    axis.title.x = element_blank(),
    panel.grid.minor = element_blank()
  ) +
  ylab("Home Value (US$)") +
  ggtitle("Housing Market in New York (1975 - 2013)")
```






## 12 Facets


Let's plot the northeast data again

```
ggplot(northeast, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line()
```



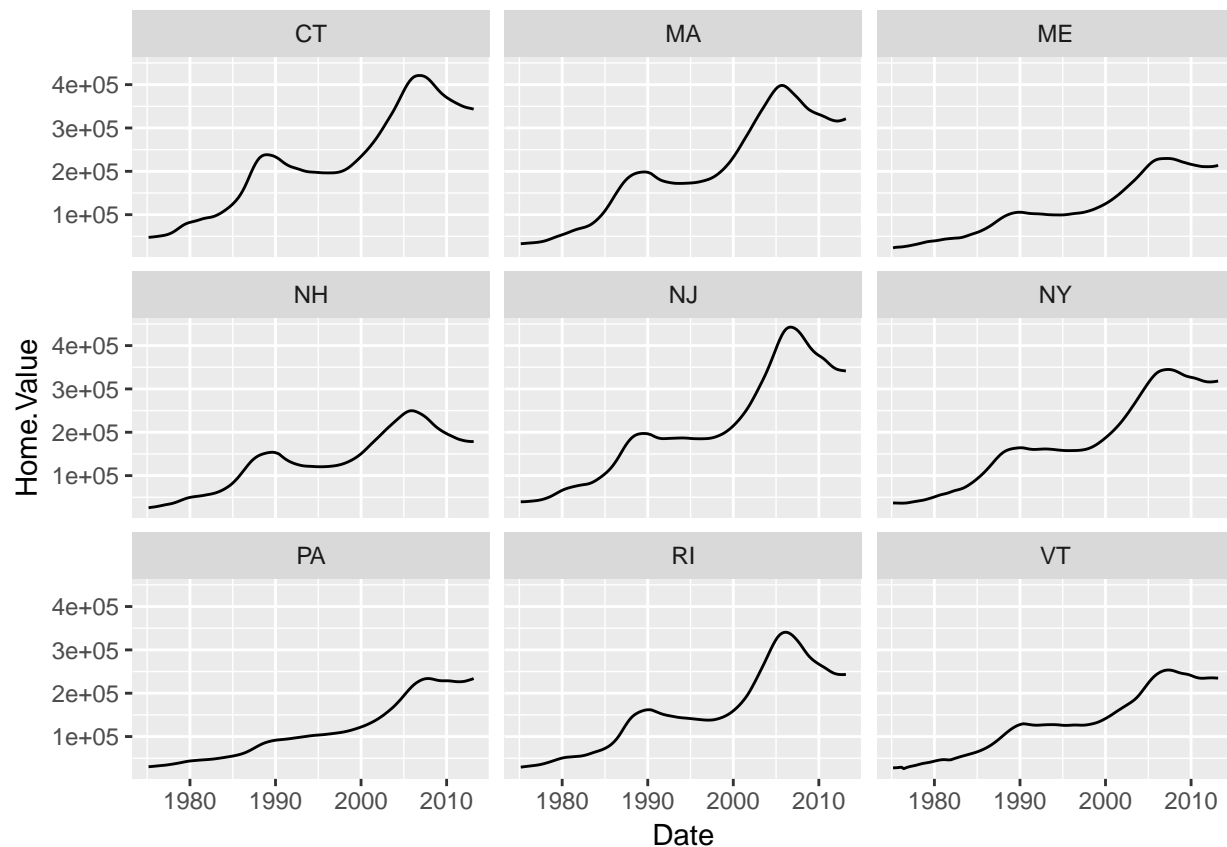
But what if we were to plot the entire dataset?

```
ggplot(housing, aes(x = Date, y = Home.Value, color = State)) +  
  geom_line()
```



The plot is not very informative anymore. We can use facets to split the plot based on the **State**

```
ggplot(northeast, aes(x = Date, y = Home.Value)) +
  geom_line() +
  facet_wrap(~State, ncol = 3)
```



## 13 Challenge


### 13.1 Recreating the Economist Graph

```
econ <- read.csv("https://raw.githubusercontent.com/altaf-ali/ggplot_tutorial/master/data/economist.csv")
head(econ)
```

##	X	Country	HDI.Rank	HDI	CPI	Region
## 1	1	Afghanistan	172	0.398	1.5	Asia Pacific
## 2	2	Albania	70	0.739	3.1	East EU Cemt Asia
## 3	3	Algeria	96	0.698	2.9	MENA
## 4	4	Angola	148	0.486	2.0	SSA
## 5	5	Argentina	45	0.797	3.0	Americas
## 6	6	Armenia	86	0.716	2.6	East EU Cemt Asia


1. Create a scatter plot of the economist data with CPI on the x-axis and HDI on the y-axis

## Corruption and human development




Sources: Transparency International; UN Human Development Report


Figure 3:




2. Color the points based on Region using hollow points



3. Add a trend line



4. The trend line is too thick compared to the circles so we need to adjust it appropriately




5. Add text labels to the points


HINT: Create a subset of countries to label since we don't want to label every point

```
target_countries <- c(
  "Russia", "Venezuela", "Iraq", "Myanmar", "Sudan",
  "Afghanistan", "Congo", "Greece", "Argentina", "Brazil",
  "India", "Italy", "China", "South Africa", "Spane",
  "Botswana", "Cape Verde", "Bhutan", "Rwanda", "France",
  "United States", "Germany", "Britain", "Barbados", "Norway", "Japan",
  "New Zealand", "Singapore"
)
labeled_countries <- subset(econ, Country %in% target_countries)
```






6. Adjust the x and y scales and use Color Brewer palette **Paired**.



7. Remove vertical grid lines and move the legend



8. Add title “Corruption and Human development”

## Corruption and Human development

